

```
In [367... import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.under_sampling import RandomUnderSampler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score, f1_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import learning_curve
```

```
In [368... magic_data = pd.read_csv('magic04.data', header=None)
magic_data.head()
```

```
Out[368...      0      1      2      3      4      5      6      7      8
0  28.7967  16.0021  2.6449  0.3918  0.1982  27.7004  22.0110  -8.2027  40.0920  81.88
1  31.6036  11.7235  2.5185  0.5303  0.3773  26.2722  23.8238  -9.9574   6.3609  205.26
2  162.0520 136.0310  4.0612  0.0374  0.0187  116.7410 -64.8580 -45.2160  76.9600  256.78
3   23.8172   9.5728  2.3385  0.6147  0.3922  27.2107  -6.4633  -7.1513  10.4490  116.73
4   75.1362  30.9205  3.1611  0.3168  0.1832  -5.5277  28.5525  21.8393   4.6480  356.46
```

```
In [369... column_names = [
    'fLength', 'fWidth', 'fSize', 'fConc', 'fConc1',
    'fAsym', 'fM3Long', 'fM3Trans', 'fAlpha', 'fDist', 'class'
]
magic_data.columns = column_names
magic_data.head()
```

```
Out[369...      fLength  fWidth  fSize  fConc  fConc1  fAsym  fM3Long  fM3Trans  fAlpha
0  28.7967  16.0021  2.6449  0.3918  0.1982  27.7004  22.0110  -8.2027  40.0920  81.
1  31.6036  11.7235  2.5185  0.5303  0.3773  26.2722  23.8238  -9.9574   6.3609  205.
2  162.0520 136.0310  4.0612  0.0374  0.0187  116.7410 -64.8580 -45.2160  76.9600  256.
3   23.8172   9.5728  2.3385  0.6147  0.3922  27.2107  -6.4633  -7.1513  10.4490  116.
4   75.1362  30.9205  3.1611  0.3168  0.1832  -5.5277  28.5525  21.8393   4.6480  356.
```

```
In [370... X = magic_data.drop('class', axis=1)
y = magic_data['class']
```

```
In [371... X_g = X[y == 'g']
y_g = y[y == 'g']
```

```
X_h = X[y == 'h']
y_h = y[y == 'h']
```

```
In [372... n_samples_h = len(X_h)
```

```
In [373... rus = RandomUnderSampler(sampling_strategy={
    'g': n_samples_h,
    'h': n_samples_h
}, random_state=42)
```


```
In [374... X_resampled, y_resampled = rus.fit_resample(X, y)
print(f"Original dataset shape: {X.shape}")
print(f"Resampled dataset shape: {X_resampled.shape}")
print(f"Original class distribution: {y.value_counts()}")
print(f"Resampled class distribution: {y_resampled.value_counts()}")
```

```
Original dataset shape: (19020, 10)
Resampled dataset shape: (13376, 10)
Original class distribution: class
g    12332
h     6688
Name: count, dtype: int64
Resampled class distribution: class
g     6688
h     6688
Name: count, dtype: int64
```


```
In [375... plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.pie(y.value_counts(), labels=y.value_counts().index, autopct='%1.1f%%')
plt.title('Original Class Distribution')
plt.subplot(1, 2, 2)
plt.pie(y_resampled.value_counts(), labels=y_resampled.value_counts().index, autopct='%1.1f%%')
plt.title('Resampled Class Distribution')

plt.tight_layout()
plt.show()
```

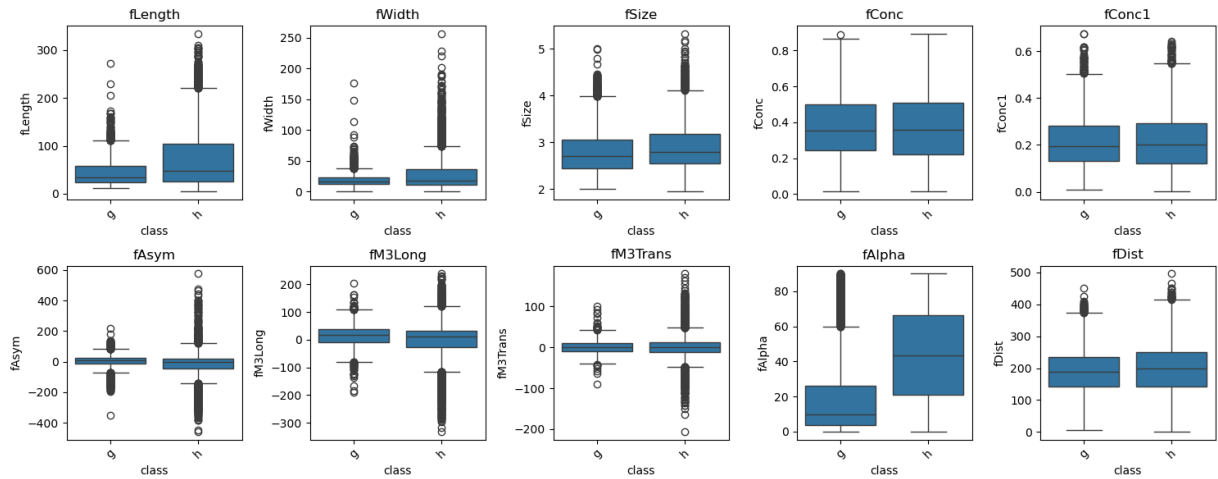
Original Class Distribution




Resampled Class Distribution



```
In [376... plt.figure(figsize=(15, 6))
for i, feature in enumerate(X_resampled.columns):
    plt.subplot(2, 5, i+1)
    sns.boxplot(x=y_resampled, y=X_resampled[feature])
    plt.title(feature)
    plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```




```
In [377... plt.figure(figsize=(12, 10))
correlation_matrix = X_resampled.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Feature Correlations')
plt.show()
```



```
In [378... feature_variance = X_resampled.var()

plt.figure(figsize=(8, 10))
feature_variance.sort_values(ascending=False).plot(kind='bar')
plt.title('Feature Variance')
plt.ylabel('Variance')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



In [379...

```
X_resampled_g = X_resampled[y_resampled == 'g']  
y_resampled_g = y_resampled[y_resampled == 'g']  
  
X_resampled_h = X_resampled[y_resampled == 'h']  
y_resampled_h = y_resampled[y_resampled == 'h']
```

In [380...

```
X_train_val_g, X_test_g, y_train_val_g, y_test_g = train_test_split(X_resampled_g,  
X_train_g, X_val_g, y_train_g, y_val_g = train_test_split(X_train_val_g, y_train_val_g,
```

```
In [381... X_train_val_h, X_test_h, y_train_val_h, y_test_h = train_test_split(X_resampled_h,
X_train_h, X_val_h, y_train_h, y_val_h = train_test_split(X_train_val_h, y_train_val_h,
```

```
In [382... X_train_scaled = np.concatenate((X_train_g, X_train_h), axis=0)
X_val_scaled = np.concatenate((X_val_g, X_val_h), axis=0)
X_test_scaled = np.concatenate((X_test_g, X_test_h), axis=0)

y_train = np.concatenate((y_train_g, y_train_h), axis=0)
y_val = np.concatenate((y_val_g, y_val_h), axis=0)
y_test = np.concatenate((y_test_g, y_test_h), axis=0)
```

```
In [383... scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_scaled)
X_val_scaled = scaler.transform(X_val_scaled)
X_test_scaled = scaler.transform(X_test_scaled)
```

```
In [384... param_grid = {
    'n_neighbors': [1, 3, 5, 7, 9, 11, 13, 15, 17],
    'metric': ['euclidean', 'manhattan'],
}

knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid=param_grid, cv=5)
grid_search.fit(X_train_scaled, y_train)
val_predictions = grid_search.predict(X_val_scaled)

print("Best parameters:", grid_search.best_params_)
print("Best validation score:", grid_search.score(X_val_scaled, y_val))
print("Accuracy:", accuracy_score(y_val, val_predictions))

best_k = grid_search.best_params_['n_neighbors']
```

Best parameters: {'metric': 'manhattan', 'n\_neighbors': 9}

Best validation score: 0.8099102947458351

Accuracy: 0.8099102947458351

```
In [386... best_k
```


```
Out[386... 9
```

```
In [395... k_values = param_grid['n_neighbors']
cv_results = grid_search.cv_results_

plt.figure(figsize=(10, 6))
for metric in param_grid['metric']:
    mask = cv_results['param_metric'] == metric
    plt.plot(k_values,
             cv_results['mean_test_score'][mask],
             marker='o',
             label=f'metric={metric}')

plt.xlabel('k (n_neighbors)')
plt.ylabel('Cross-validation Accuracy')
plt.title('KNN Performance vs k Value')
plt.legend()
```

```
plt.grid(True)  
plt.show()
```



In [388... `help(KNeighborsClassifier)`

Help on class KNeighborsClassifier in module sklearn.neighbors.\_classification:

```
class KNeighborsClassifier(sklearn.neighbors._base.KNeighborsMixin, sklearn.base.ClassifierMixin, sklearn.neighbors._base.NeighborsBase)
```

```
| KNeighborsClassifier(n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=30, p=2, metric='minkowski', metric_params=None, n_jobs=None)
```

```
|
```

```
| Classifier implementing the k-nearest neighbors vote.
```

```
|
```

```
| Read more in the :ref:`User Guide <classification>`.
```

```
|
```

```
| Parameters
```

```
| -----
```

```
| n_neighbors : int, default=5
```

```
|     Number of neighbors to use by default for :meth:`kneighbors` queries.
```

```
|
```

```
| weights : {'uniform', 'distance'}, callable or None, default='uniform'
```

```
|     Weight function used in prediction. Possible values:
```

```
|
```

```
| - 'uniform' : uniform weights. All points in each neighborhood  
| are weighted equally.
```

```
| - 'distance' : weight points by the inverse of their distance.  
| in this case, closer neighbors of a query point will have a  
| greater influence than neighbors which are further away.
```

```
| - [callable] : a user-defined function which accepts an  
| array of distances, and returns an array of the same shape  
| containing the weights.
```

```
|
```

```
| Refer to the example entitled
```

```
| :ref:`sphx_glr_auto_examples_neighbors_plot_classification.py`
```

```
| showing the impact of the `weights` parameter on the decision  
| boundary.
```

```
|
```

```
| algorithm : {'auto', 'ball_tree', 'kd_tree', 'brute'}, default='auto'
```

```
|     Algorithm used to compute the nearest neighbors:
```

```
|
```

```
| - 'ball_tree' will use :class:`BallTree`
```

```
| - 'kd_tree' will use :class:`KDTree`
```

```
| - 'brute' will use a brute-force search.
```

```
| - 'auto' will attempt to decide the most appropriate algorithm  
| based on the values passed to :meth:`fit` method.
```

```
|
```

```
| Note: fitting on sparse input will override the setting of  
| this parameter, using brute force.
```

```
|
```

```
| leaf_size : int, default=30
```

```
|     Leaf size passed to BallTree or KDTree. This can affect the  
| speed of the construction and query, as well as the memory  
| required to store the tree. The optimal value depends on the  
| nature of the problem.
```

```
|
```

```
| p : float, default=2
```

```
|     Power parameter for the Minkowski metric. When p = 1, this is equivalent  
| to using manhattan_distance (l1), and euclidean_distance (l2) for p = 2.
```

```
| For arbitrary p, minkowski_distance (l_p) is used. This parameter is expected
```

```
|
```



to be positive.

`metric` : str or callable, default='minkowski'

Metric to use for distance computation. Default is "minkowski", which results in the standard Euclidean distance when  $p = 2$ . See the documentation of `scipy.spatial.distance` <https://docs.scipy.org/doc/scipy/reference/spatial.distance.html> and the metrics listed in `class:~sklearn.metrics.pairwise.distance_metrics` for valid metric values.

If `metric` is "precomputed", `X` is assumed to be a distance matrix and must be square during fit. `X` may be a `:term:`sparse graph``, in which case only "nonzero" elements may be considered neighbors.

If `metric` is a callable function, it takes two arrays representing 1D vectors as inputs and must return one value indicating the distance between those vectors. This works for Scipy's metrics, but is less efficient than passing the metric name as a string.

`metric_params` : dict, default=None

Additional keyword arguments for the metric function.

`n_jobs` : int, default=None

The number of parallel jobs to run for neighbors search.

`None` means 1 unless in a `:obj:`joblib.parallel_backend`` context.

`-1` means using all processors. See `:term:`Glossary <n_jobs>`` for more details.

Doesn't affect `:meth:`fit`` method.

Attributes

-----

`classes_` : array of shape (n\_classes,)

Class labels known to the classifier

`effective_metric_` : str or callable

The distance metric used. It will be same as the ``metric`` parameter or a synonym of it, e.g. 'euclidean' if the ``metric`` parameter set to 'minkowski' and ``p`` parameter set to 2.

`effective_metric_params_` : dict

Additional keyword arguments for the metric function. For most metrics will be same with ``metric_params`` parameter, but may also contain the ``p`` parameter value if the ``effective_metric_`` attribute is set to 'minkowski'.

`n_features_in_` : int

Number of features seen during `:term:`fit``.

.. versionadded:: 0.24

`feature_names_in_` : ndarray of shape (``n_features_in_``),

Names of features seen during `:term:`fit``. Defined only when ``X`` has feature names that are all strings.

.. versionadded:: 1.0

```

n_samples_fit_ : int
    Number of samples in the fitted data.

outputs_2d_ : bool
    False when `y`'s shape is (n_samples, ) or (n_samples, 1) during fit
    otherwise True.

```

#### See Also

-----

RadiusNeighborsClassifier: Classifier based on neighbors within a fixed radius.  
 KNeighborsRegressor: Regression based on k-nearest neighbors.  
 RadiusNeighborsRegressor: Regression based on neighbors within a fixed radius.  
 NearestNeighbors: Unsupervised learner for implementing neighbor searches.

#### Notes

-----

See :ref:`Nearest Neighbors <neighbors>` in the online documentation  
 for a discussion of the choice of ``algorithm`` and ``leaf\_size``.

.. warning::

Regarding the Nearest Neighbors algorithms, if it is found that two  
 neighbors, neighbor `k+1` and `k`, have identical distances  
 but different labels, the results will depend on the ordering of the  
 training data.

[https://en.wikipedia.org/wiki/K-nearest\\_neighbor\\_algorithm](https://en.wikipedia.org/wiki/K-nearest_neighbor_algorithm)

#### Examples

-----

```

>>> X = [[0], [1], [2], [3]]
>>> y = [0, 0, 1, 1]
>>> from sklearn.neighbors import KNeighborsClassifier
>>> neigh = KNeighborsClassifier(n_neighbors=3)
>>> neigh.fit(X, y)
KNeighborsClassifier(...)
>>> print(neigh.predict([[1.1]]))
[0]
>>> print(neigh.predict_proba([[0.9]]))
[[0.666... 0.333...]]

```

#### Method resolution order:

```

KNeighborsClassifier
sklearn.neighbors._base.KNeighborsMixin
sklearn.base.ClassifierMixin
sklearn.neighbors._base.NeighborsBase
sklearn.base.MultiOutputMixin
sklearn.base.BaseEstimator
sklearn.utils._estimator_html_repr._HTMLDocumentationLinkMixin
sklearn.utils._metadata_requests._MetadataRequester
builtins.object

```

#### Methods defined here:

```

__init__(self, n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=

```

```

30, p=2, metric='minkowski', metric_params=None, n_jobs=None)
|     Initialize self. See help(type(self)) for accurate signature.
|
|     fit(self, X, y)
|         Fit the k-nearest neighbors classifier from the training dataset.
|
|         Parameters
|         -----
|         X : {array-like, sparse matrix} of shape (n_samples, n_features) or
(n_samples, n_samples) if metric='precomputed'
|             Training data.
|
|         y : {array-like, sparse matrix} of shape (n_samples,) or           (n_
samples, n_outputs)
|             Target values.
|
|         Returns
|         -----
|         self : KNeighborsClassifier
|             The fitted k-nearest neighbors classifier.
|
|     predict(self, X)
|         Predict the class labels for the provided data.
|
|         Parameters
|         -----
|         X : {array-like, sparse matrix} of shape (n_queries, n_features),
or (n_queries, n_indexed) if metric == 'precomputed'
|             Test samples.
|
|         Returns
|         -----
|         y : ndarray of shape (n_queries,) or (n_queries, n_outputs)
|             Class labels for each data sample.
|
|     predict_proba(self, X)
|         Return probability estimates for the test data X.
|
|         Parameters
|         -----
|         X : {array-like, sparse matrix} of shape (n_queries, n_features),
or (n_queries, n_indexed) if metric == 'precomputed'
|             Test samples.
|
|         Returns
|         -----
|         p : ndarray of shape (n_queries, n_classes), or a list of n_outputs
of such arrays if n_outputs > 1.
|             The class probabilities of the input samples. Classes are ordered
|             by lexicographic order.
|
|     set_score_request(self: sklearn.neighbors._classification.KNeighborsClassifier,
*, sample_weight: Union[bool, NoneType, str] = '$UNCHANGED$') -> sklearn.neighbors._
classification.KNeighborsClassifier from sklearn.utils._metadata_requests.RequestMet
hod.__get__.<locals>
|         Request metadata passed to the ``score`` method.

```

Note that this method is only relevant if `enable_metadata_routing=True` (see `:func:sklearn.set_config`). Please see `:ref:User Guide <metadata_routing>` on how the routing mechanism works.

The options for each parameter are:

- `True`: metadata is requested, and passed to `score` if provided. The request is ignored if metadata is not provided.
- `False`: metadata is not requested and the meta-estimator will not pass it to `score`.
- `None`: metadata is not requested, and the meta-estimator will raise an error if the user provides it.
- `str`: metadata should be passed to the meta-estimator with this given a `alias` instead of the original name.

The default (`sklearn.utils.metadata_routing.UNCHANGED`) retains the existing request. This allows you to change the request for some parameters and not others.

.. versionadded:: 1.3

.. note::

This method is only relevant if this estimator is used as a sub-estimator of a meta-estimator, e.g. used inside a `sklearn.pipeline.Pipeline`. Otherwise it has no effect.

Parameters

-----

`sample_weight` : str, True, False, or None, default=`sklearn.utils.metadata_routing.UNCHANGED`  
Metadata routing for `sample_weight` parameter in `score`.

Returns

-----

`self` : object  
The updated object.

-----  
Data and other attributes defined here:

`__abstractmethods__` = frozenset()

`__annotations__` = {'\_parameter\_constraints': <class 'dict'>}

`__slotnames__` = []

-----  
Methods inherited from `sklearn.neighbors._base.KNeighborsMixin`:

`kneighbors(self, X=None, n_neighbors=None, return_distance=True)`  
Find the K-neighbors of a point.

Returns indices of and distances to the neighbors of each point.

#### Parameters

-----

X : {array-like, sparse matrix}, shape (n\_queries, n\_features),  
or (n\_queries, n\_indexed) if metric == 'precomputed', default=None  
The query point or points.  
If not provided, neighbors of each indexed point are returned.  
In this case, the query point is not considered its own neighbor.

n\_neighbors : int, default=None

Number of neighbors required for each sample. The default is the value passed to the constructor.

return\_distance : bool, default=True

Whether or not to return the distances.

#### Returns

-----

neigh\_dist : ndarray of shape (n\_queries, n\_neighbors)

Array representing the lengths to points, only present if return\_distance=True.

neigh\_ind : ndarray of shape (n\_queries, n\_neighbors)

Indices of the nearest points in the population matrix.

#### Examples

-----

In the following example, we construct a NearestNeighbors class from an array representing our data set and ask who's the closest point to [1,1,1]

```
>>> samples = [[0., 0., 0.], [0., .5, 0.], [1., 1., .5]]
>>> from sklearn.neighbors import NearestNeighbors
>>> neigh = NearestNeighbors(n_neighbors=1)
>>> neigh.fit(samples)
NearestNeighbors(n_neighbors=1)
>>> print(neigh.kneighbors([[1., 1., 1.]])
(array([[0.5]]), array([[2]]))
```

As you can see, it returns [[0.5]], and [[2]], which means that the element is at distance 0.5 and is the third element of samples (indexes start at 0). You can also query for multiple points:

```
>>> X = [[0., 1., 0.], [1., 0., 1.]]
>>> neigh.kneighbors(X, return_distance=False)
array([[1],
       [2]]...)
```

kneighbors\_graph(self, X=None, n\_neighbors=None, mode='connectivity')  
Compute the (weighted) graph of k-Neighbors for points in X.

#### Parameters

-----

X : {array-like, sparse matrix} of shape (n\_queries, n\_features),

```

or (n_queries, n_indexed) if metric == 'precomputed', default=None
    The query point or points.
    If not provided, neighbors of each indexed point are returned.
    In this case, the query point is not considered its own neighbor.
    For ``metric='precomputed'`` the shape should be
    (n_queries, n_indexed). Otherwise the shape should be
    (n_queries, n_features).

n_neighbors : int, default=None
    Number of neighbors for each sample. The default is the value
    passed to the constructor.

mode : {'connectivity', 'distance'}, default='connectivity'
    Type of returned matrix: 'connectivity' will return the
    connectivity matrix with ones and zeros, in 'distance' the
    edges are distances between points, type of distance
    depends on the selected metric parameter in
    NearestNeighbors class.

Returns
-----
A : sparse-matrix of shape (n_queries, n_samples_fit)
    `n_samples_fit` is the number of samples in the fitted data.
    `A[i, j]` gives the weight of the edge connecting `i` to `j`.
    The matrix is of CSR format.

See Also
-----
NearestNeighbors.radius_neighbors_graph : Compute the (weighted) graph
    of Neighbors for points in X.

Examples
-----
>>> X = [[0], [3], [1]]
>>> from sklearn.neighbors import NearestNeighbors
>>> neigh = NearestNeighbors(n_neighbors=2)
>>> neigh.fit(X)
NearestNeighbors(n_neighbors=2)
>>> A = neigh.kneighbors_graph(X)
>>> A.toarray()
array([[1., 0., 1.],
       [0., 1., 1.],
       [1., 0., 1.]])

-----
Data descriptors inherited from sklearn.neighbors._base.KNeighborsMixin:

__dict__
    dictionary for instance variables

__weakref__
    list of weak references to the object

-----
Methods inherited from sklearn.base.ClassifierMixin:

```

```

score(self, X, y, sample_weight=None)
    Return the mean accuracy on the given test data and labels.

    In multi-label classification, this is the subset accuracy
    which is a harsh metric since you require for each sample that
    each label set be correctly predicted.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Test samples.

    y : array-like of shape (n_samples,) or (n_samples, n_outputs)
        True labels for `X`.

    sample_weight : array-like of shape (n_samples,), default=None
        Sample weights.

    Returns
    -----
    score : float
        Mean accuracy of ``self.predict(X)`` w.r.t. `y`.

-----
Methods inherited from sklearn.base.BaseEstimator:

__getstate__(self)
    Helper for pickle.

__repr__(self, N_CHAR_MAX=700)
    Return repr(self).

__setstate__(self, state)

__sklearn_clone__(self)

get_params(self, deep=True)
    Get parameters for this estimator.

    Parameters
    -----
    deep : bool, default=True
        If True, will return the parameters for this estimator and
        contained subobjects that are estimators.

    Returns
    -----
    params : dict
        Parameter names mapped to their values.

set_params(self, **params)
    Set the parameters of this estimator.

    The method works on simple estimators as well as on nested objects
    (such as :class:`~sklearn.pipeline.Pipeline`). The latter have
    parameters of the form ``<component>__<parameter>`` so that it's

```

```

    possible to update each component of a nested object.

    Parameters
    -----
    **params : dict
        Estimator parameters.

    Returns
    -----
    self : estimator instance
        Estimator instance.

    -----
    Methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

    get_metadata_routing(self)
        Get metadata routing of this object.

        Please check :ref:`User Guide <metadata_routing>` on how the routing
        mechanism works.

    Returns
    -----
    routing : MetadataRequest
        A :class:`~sklearn.utils.metadata_routing.MetadataRequest` encapsulating
        routing information.

    -----
    Class methods inherited from sklearn.utils._metadata_requests._MetadataRequester:

r:
    __init_subclass__(**kwargs)
        Set the ``set_{method}_request`` methods.

        This uses PEP-487 [1]_ to set the ``set_{method}_request`` methods. It
        looks for the information available in the set default values which are
        set using ``__metadata_request__`` class attributes, or inferred
        from method signatures.

        The ``__metadata_request__`` class attributes are used when a method
        does not explicitly accept a metadata through its arguments or if the
        developer would like to specify a request value for those metadata
        which are different from the default ``None``.

    References
    -----
    .. [1] https://www.python.org/dev/peps/pep-0487

```

In [389...

```

best_knn = KNeighborsClassifier(n_neighbors=best_k, metric=best_params['metric'])
best_knn.fit(X_train_scaled, y_train)

```



Out[389...

KNeighborsClassifier

KNeighborsClassifier(metric='manhattan', n\_neighbors=9)


In [390...

```

train_sizes, train_scores, val_scores = learning_curve(
    best_knn, X_train_scaled, y_train,
    train_sizes=np.linspace(0.1, 1.0, 10),
    cv=5, n_jobs=-1)

plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_scores.mean(axis=1), label='Training score')
plt.plot(train_sizes, val_scores.mean(axis=1), label='Cross-validation score')
plt.xlabel('Training Examples')
plt.ylabel('Score')
plt.title('Learning Curves')
plt.legend(loc='best')
plt.grid(True)
plt.show()

```



In [391...

```
y_test_pred = best_knn.predict(X_test_scaled)
```

In [392...

```

accuracy = accuracy_score(y_test, y_test_pred)
precision = precision_score(y_test, y_test_pred, average='weighted')
recall = recall_score(y_test, y_test_pred, average='weighted')
f1 = f1_score(y_test, y_test_pred, average='weighted')
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1-score: {f1}")

```

Accuracy: 0.8074240159441953  
Precision: 0.8190425520951419  
Recall: 0.8074240159441953  
F1-score: 0.8056546528713431

```
In [393... cm = confusion_matrix(y_test, y_test_pred)

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

